



Research Question: How can we train a general-purpose Reasoning LLM across diverse domains via RL?


Results Comparison



Sequential RL (Sec. 3.1)



Mixed RL (Sec. 3.3)



Modular Gradient Surgery (Ours, Sec. 4&5)

